

强化学习的数学原理

👤 Xuhang Ye

✉️ cquleaf@yexuhang.com

创建于: 2025 年 11 月 11 日

更新于: 2025 年 11 月 21 日

目录

1 Basic Concepts	3
1.1 Basic concepts in reinforcement learning	3
1.2 Markov decision process (MDP)	4
2 Bellman Equation	5
2.1 Motivating examples	5
2.2 State value	6
2.3 Bellman equation: Derivation	8
2.4 Bellman equation: Matrix-vector form	9
3 插图示例	10
4 彩色信息框示例	11



1 Basic Concepts

1.1 Basic concepts in reinforcement learning

1. ***State***: The status of the agent with respect to the environment.
2. ***State space***: The set of all states $\mathcal{S} = \{s_i\}_{i=1}^n$.
3. ***Action***: For each state, there are several possible actions: a_1, \dots, a_n .
4. ***Action space of a state***: The set of all possible actions of a state $\mathcal{A}(s_i) = \{a_i\}_{i=1}^n$.
5. ***State transition***: When taking an action, the agent may move from one state to another.
6. ***State transition probability***: Use (conditional) probability to describe state transition!
7. ***Policy***: It tells the agent what actions to take at a state.
8. ***Reward***: A real number we get after taking an action. (Usually, positive reward represents encouragement and negative reward represents punishment)
9. ***Trajectory***: A state-action-reward chain.
10. ***Return***: The return of one trajectory is the sum of all the rewards collected along the trajectory.
11. ***Discounted rate***: $\gamma \in [0, 1]$, its roles are making the sum become finite and balancing the far and near future rewards: $\gamma \rightarrow 0$ represents the near future and $\gamma \rightarrow 1$ represents the far future.
12. ***Episode***: When interacting with the environment following a policy, the agent may stop at some *terminal states*. The resulting trajectory is called an episode (or a trial). An episode is usually assumed to be a finite trajectory. Tasks with episodes are called *episode tasks*.

1.2 Markov decision process (MDP)

1. Sets:

- **State:** The set of state \mathcal{S}
- **Action:** The set of actions $\mathcal{A}(s)$ is associated for state $s \in \mathcal{S}$.
- **Reward:** The set of rewards $\mathcal{R}(s, a)$.

2. Probability distribution:

- **State transition probability:** At state s , taking action a , the probability to transit to state s' is $p(s' | s, a)$.
- **Reward probability:** At state s , taking action a , the probability to get reward r is $p(r | s, a)$.

3. Policy:

At state s , the probability to choose action a is $\pi(a | s)$.

4. Markov property:

memoryless property

$$\begin{aligned} p(s_{t+1} | a_{t+1}, s_t, \dots, a_1, s_0) &= p(s_{t+1} | a_{t+1}, s_t), \\ p(r_{t+1} | a_{t+1}, s_t, \dots, a_1, s_0) &= p(r_{t+1} | a_{t+1}, s_t). \end{aligned} \tag{1}$$

Markov decision process becomes Markov process once the policy is given!

2 Bellman Equation

2.1 Motivating examples

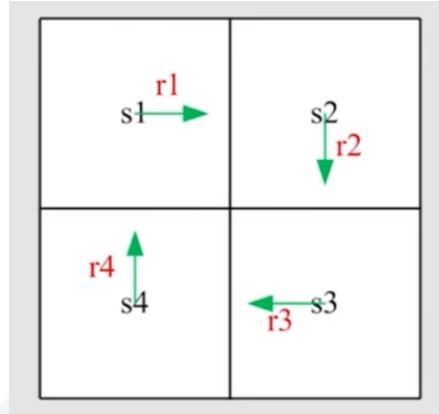


图 1: The example to introduce bellman equation

How to calculate return?

Method 1: by definition

Let v_i denote the return obtained starting from s_i ($i = 1, 2, 3, 4$)

$$\begin{aligned} v_1 &= r_1 + \gamma r_2 + \gamma^2 r_3 + \dots \\ v_2 &= r_2 + \gamma r_3 + \gamma^2 r_4 + \dots \\ v_3 &= r_3 + \gamma r_4 + \gamma^2 r_1 + \dots \\ v_4 &= r_4 + \gamma r_1 + \gamma^2 r_2 + \dots \end{aligned} \tag{2}$$

Method 2: bootstrapping

$$\begin{aligned} v_1 &= r_1 + \gamma (r_2 + \gamma r_3 + \dots) = r_1 + \gamma v_2 \\ v_2 &= r_2 + \gamma (r_3 + \gamma r_4 + \dots) = r_2 + \gamma v_3 \\ v_3 &= r_3 + \gamma (r_4 + \gamma r_1 + \dots) = r_3 + \gamma v_4 \\ v_4 &= r_4 + \gamma (r_1 + \gamma r_2 + \dots) = r_4 + \gamma v_1 \end{aligned} \tag{3}$$

How to solve the above equations?

Write in the following matrix-vector form:

$$\underbrace{\begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ v_4 \end{bmatrix}}_{\mathbf{v}} = \underbrace{\begin{bmatrix} r_1 \\ r_2 \\ r_3 \\ r_4 \end{bmatrix}}_{\mathbf{r}} + \underbrace{\begin{bmatrix} \gamma v_2 \\ \gamma v_3 \\ \gamma v_4 \\ \gamma v_1 \end{bmatrix}}_{\mathbf{P}\mathbf{v}} = \underbrace{\begin{bmatrix} r_1 \\ r_2 \\ r_3 \\ r_4 \end{bmatrix}}_{\mathbf{r}} + \gamma \underbrace{\begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \end{bmatrix}}_{\mathbf{P}} \underbrace{\begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ v_4 \end{bmatrix}}_{\mathbf{v}} \quad (4)$$

which can be rewritten as

$$\mathbf{v} = \mathbf{r} + \gamma \mathbf{Pv} \quad (5)$$

This is **Bellman equation** (for this specific deterministic problem)

1. Though simple, it demonstrates the core idea: the value of one state relies on the values of other states.
2. A matrix-vector form is more clear to see how to solve the state values.

2.2 State value

Consider the following single-step process:

$$S_t \xrightarrow{A_t} R_{t+1}, S_{t+1} \quad (6)$$

1. $t, t+1$: discrete time instances
2. S_t : **state** at time t
3. A_t : the **action** taken at state S_t
4. R_{t+1} : the **reward** obtained after taking A_t
5. S_{t+1} : the **state transited** to after taking A_t

Note that S_t, A_t, R_{t+1} are all *random variables*.

This step is governed by the following probability distributions:

1. $S_t \rightarrow A_t$ is governed by $\pi(A_t = a | S_t = s)$
2. $S_t, A_t \rightarrow R_{t+1}$ is governed by $p(R_{t+1} = r | S_t = s, A_t = a)$
3. $S_t, A_t \rightarrow S_{t+1}$ is governed by $p(S_{t+1} = s' | S_t = s, A_t = a)$

At this moment, we assume we know the model (i.e., the probability distributions)! Consider the following multi-step trajectory:

$$S_t \xrightarrow{A_t} R_{t+1}, S_{t+1} \xrightarrow{A_{t+1}} R_{t+2}, S_{t+2} \xrightarrow{A_{t+2}} R_{t+3}, \dots \quad (7)$$

The discounted return is

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \quad (8)$$

Where:

- $\gamma \in [0, 1]$ is a discounted rate.
- G_t is also a random variable since R_{t+1}, R_{t+2}, \dots are random variables.

The expectation (or called expected value or mean) of G_t is defined as the *state-value function* or simply *state value*:

$$v_\pi(s) = \mathbb{E}[G_t | S_t = s] \quad (9)$$

Remarks:

1. It is a function of s . It is a conditional expectation with the condition that the state starts from s .
2. It is based on the policy π . For a different policy, the state value may be different.
3. It represents the “value” of a state. If the state value is greater, then the policy is better because greater cumulative rewards can be obtained.

Q: What is the relationship between **return** and **state value**?

A: The state value is the mean of all possible returns that can be obtained starting from a state. If everything — $\pi(a | s)$, $p(r | s, a)$, $p(s' | s, a)$ —is deterministic, then state value is the same as return.

2.3 Bellman equation: Derivation

Consider a random trajectory:

$$S_t \xrightarrow{A_t} R_{t+1}, S_{t+1} \xrightarrow{A_{t+1}} R_{t+2}, S_{t+2} \xrightarrow{A_{t+2}} R_{t+3}, \dots \quad (10)$$

The return G_t can be written as

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots, \\ &= R_{t+1} + \gamma (R_{t+2} + \gamma R_{t+3} + \dots), \\ &= R_{t+1} + \gamma G_{t+1}, \end{aligned} \quad (11)$$

Then, it follows from the definition of the state value that

$$\begin{aligned} v_\pi(s) &= \mathbb{E}[G_t | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \mathbb{E}[R_{t+1} | S_t = s] + \gamma \mathbb{E}[G_{t+1} | S_t = s] \end{aligned} \quad (12)$$

Next, calculate the two terms, respectively.

First, calculate the first term $\mathbb{E}[R_{t+1} | S_t = s]$:

$$\begin{aligned} \mathbb{E}[R_{t+1} | S_t = s] &= \sum_a \pi(a | s) \mathbb{E}[R_{t+1} | S_t = s, A_t = a] \\ &= \sum_a \pi(a | s) \sum_r p(r | s, a) r \end{aligned} \quad (13)$$

Note: This is the mean of *immediate rewards*

Second, calculate the second term $\mathbb{E}[G_{t+1} | S_t = s]$

$$\begin{aligned} \mathbb{E}[G_{t+1} | S_t = s] &= \sum_{s'} \mathbb{E}[G_{t+1} | S_t = s, S_{t+1} = s'] p(s' | s) \\ &= \sum_{s'} \mathbb{E}[G_{t+1} | S_{t+1} = s'] p(s' | s) \\ &= \sum_{s'} v_\pi(s') p(s' | s) \\ &= \sum_{s'} v_\pi(s') \sum_a p(s' | s, a) \pi(a | s) \end{aligned} \quad (14)$$

Note: This is the mean of *future rewards*

Theorem 2.1: Bellman equation

Therefore, we have:

$$\begin{aligned}
 v_{\pi}(s) &= \mathbb{E}[R_{t+1} | S_t = s] + \gamma \mathbb{E}[G_{t+1} | S_t = s], \\
 &= \underbrace{\sum_a \pi(a | s) \sum_r p(r | s, a)r}_{\text{mean of immediate rewards}} + \underbrace{\gamma \sum_a \pi(a | s) \sum_{s'} p(s' | s, a) v_{\pi}(s')}_{\text{mean of future rewards}}, \\
 &= \sum_a \pi(a | s) \left[\sum_r p(r | s, a)r + \gamma \sum_{s'} p(s' | s, a) v_{\pi}(s') \right], \quad \forall s \in \mathcal{S}.
 \end{aligned} \tag{15}$$

Highlights:

1. The above equation is called **the Bellman equation**, which characterizes the relationship among the state-value functions of different states.
2. It consists of two terms: the immediate reward term and the future reward term.
3. A set of equations: every state has an equation like this.(important!!!)
4. $v_{\pi}(s)$ and $v_{\pi}(s')$ are state values to be calculated. Using **bootstrapping** to solve it!
5. $\pi(a | s)$ is a given policy. Solving the equation is called policy evaluation.
6. $p(r | s, a)$ and $p(s' | s, a)$ represent the dynamic model.

2.4 Bellman equation: Matrix-vector form

3 插图示例



图 2: 示例图片



4 彩色信息框示例

Definition 4.1: 具身马尔可夫决策过程 (EMDP)

一个具身马尔可夫决策过程定义为五元组 $\langle \mathcal{S}, \mathcal{A}, T, R, \gamma \rangle$, 其中:

- $\mathcal{S} = \mathbb{R}^2 \times [0, 2\pi)$ 为状态空间 (位置 + 朝向);
- $\mathcal{A} = \{\text{前进, 左转, 右转}\}$ 为离散动作空间;
- $T(s' | s, a)$ 为状态转移函数;
- $R(s, a) = -\|s_{\text{pos}} - s_{\text{goal}}\|$ 为稀疏奖励 (负欧氏距离);
- $\gamma \in (0, 1)$ 为折扣因子。

Assumption 4.1: 可观测性

能体在每一步都能精确观测当前状态 $s_t \in \mathcal{S}$ 。

Assumption 4.2: 确定性动力学

状态转移是确定性的, 即 $s_{t+1} = f(s_t, a_t)$, 无环境噪声。

该算法的正确性由以下定理保证。

Lemma 4.1: 贝尔曼最优性

函数 V^* 是贝尔曼最优算子 \mathcal{T} 的唯一不动点, 即 $V^* = \mathcal{T}V^*$, 其中

$$\mathcal{T}V(s) := \max_{a \in \mathcal{A}} \left[R(s, a) + \gamma \sum_{s'} T(s' | s, a) V(s') \right].$$

Theorem 4.1: 值迭代收敛性

???? 下, 值迭代算法生成的序列 $\{V_k\}$ 以指数速率收敛到 V^* , 即

$$\|V_k - V^*\|_\infty \leq \frac{2\gamma^k}{1-\gamma} \|V_1 - V_0\|_\infty.$$

证明. 由于状态空间 \mathcal{S} 在假设下为有限 (或可离散化), 且 $\gamma < 1$, 贝尔曼算子 \mathcal{T} 是 γ -压缩映射。由巴拿赫不动点定理, 迭代 $V_{k+1} = \mathcal{T}V_k$ 收敛到唯一不动点 V^* , 且误差界如上所述。□

Example 4.1: 2D 网格世界导航

虑一个 5×5 的网格世界，机器人从左下角 $(0,0)$ 出发，目标为右上角 $(4,4)$ 。状态 $s = (x, y, \theta)$ ，其中 $\theta \in \{0, \pi/2, \pi, 3\pi/2\}$ 。动作 $A = \{\text{前进}, \text{左转}, \text{右转}\}$ 控制朝向与位置。在 ?? 下，值迭代可精确计算到达目标的最短路径。